

Imports

```
In [ ]: !pip install datasets
```

```
In [1]: from google.colab import drive
drive.mount('/content/drive')
```

Mounted at /content/drive

```
In [53]: import warnings
warnings.filterwarnings("ignore")
```

```
In [54]: import csv
import json
import random
import pandas as pd

import seaborn as sns
import matplotlib.pyplot as plt

from mlxtend.preprocessing import TransactionEncoder
from mlxtend.frequent_patterns import apriori, fpmax, fpgrowth
from mlxtend.frequent_patterns import association_rules

import tensorflow as tf
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Conv2D, MaxPooling2D, Flatten, Dense

from datasets import Dataset, DatasetDict
```

Read Grocery data

```
In [6]: data = []

with open("/content/drive/MyDrive/DM3/Grocery_Items_23.csv", "r") as file_:
    csv_reader = csv.reader(file_)

    next(csv_reader)

    for row in csv_reader:
        row = list(filter(lambda x: x != '', row))
        data.append(row)
```

1 (c).

```
In [7]: all_items = []
        for row in data:
            for item in row:
                all_items.append(item)
```

```
In [8]: len(set(all_items))
```

```
Out[8]: 166
```

The Number of Unique Items is 166

```
In [9]: len(data)
```

```
Out[9]: 8000
```

Total Number of records in 8000

```
In [10]: from collections import Counter

          my_counter = Counter(all_items)

          max_element = my_counter.most_common(1)[0]
          print(max_element)

          ('whole milk', 1366)
```

The most popular item is 'whole milk' and 1366 transaction contain this item

1(d)

```
In [11]: def fit_association_rules(dataset, support, confidence):
          te = TransactionEncoder()
          te_ary = te.fit(dataset).transform(dataset)
          df = pd.DataFrame(te_ary, columns=te.columns_)
          frequent_itemsets = fpgrowth(df, min_support=support, use_colnames=True)
          rules = association_rules(frequent_itemsets, metric="confidence", min_threshold=confidence, num_itemsets=len(df))
          return rules
```

```
In [12]: rules = fit_association_rules(data, 0.01, 0.08)
```

In [13]: rules

Out[13]:

	antecedents	consequents	antecedent support	consequent support	support	confidence	lift	representative
0	(other vegetables)	(whole milk)	0.124250	0.161875	0.015500	0.124748	0.770647	
1	(whole milk)	(other vegetables)	0.161875	0.124250	0.015500	0.095753	0.770647	
2	(rolls/buns)	(whole milk)	0.111250	0.161875	0.014750	0.132584	0.819053	
3	(whole milk)	(rolls/buns)	0.161875	0.111250	0.014750	0.091120	0.819053	
4	(rolls/buns)	(other vegetables)	0.111250	0.124250	0.010875	0.097753	0.786743	
5	(other vegetables)	(rolls/buns)	0.124250	0.111250	0.010875	0.087525	0.786743	
6	(yogurt)	(whole milk)	0.086500	0.161875	0.011875	0.137283	0.848082	
7	(soda)	(whole milk)	0.098250	0.161875	0.013500	0.137405	0.848831	
8	(whole milk)	(soda)	0.161875	0.098250	0.013500	0.083398	0.848831	
9	(soda)	(other vegetables)	0.098250	0.124250	0.010500	0.106870	0.860123	
10	(other vegetables)	(soda)	0.124250	0.098250	0.010500	0.084507	0.860123	


1(e)

```

In [14]: min_supports_values = [0.001, 0.005, 0.01]
min_confidence_values = [0.05, 0.075, 0.1]
heat_map = []
for min_confidence in min_confidence_values:
    temp = []
    for min_support in min_supports_values:
        rules = fit_association_rules(data,min_support,min_confidence)
        temp.append(len(rules))
    heat_map.append(temp)

```

```
In [15]: sns.heatmap(heat_map, annot=True, fmt='d', cmap='viridis')
plt.xticks(ticks=[0.5, 1.5, 2.5], labels=min_supports_values)
plt.yticks(ticks=[0.5, 1.5, 2.5], labels=min_confidence_values)
plt.xlabel("msv")
plt.ylabel("mct")
plt.show()
```



2 Create and compile model

```
In [17]: model = Sequential([
    Conv2D(8, (3, 3), activation='relu', input_shape=(256, 256, 3)),
    MaxPooling2D((2, 2)),
    Conv2D(4, (3, 3), activation='relu'),
    MaxPooling2D((2, 2)),
    Flatten(),
    Dense(8, activation='relu'),
    Dense(4, activation='softmax')
])
```

```
In [18]: model.compile(optimizer='adam',
    loss='categorical_crossentropy',
    metrics=[tf.keras.metrics.CategoricalAccuracy(name='accuracy')])
```

In [19]: `model.summary()`

Model: "sequential"

Layer (type)	Output Shape
conv2d (Conv2D)	(None, 254, 254, 8)
max_pooling2d (MaxPooling2D)	(None, 127, 127, 8)
conv2d_1 (Conv2D)	(None, 125, 125, 4)
max_pooling2d_1 (MaxPooling2D)	(None, 62, 62, 4)
flatten (Flatten)	(None, 15376)
dense (Dense)	(None, 8)
dense_1 (Dense)	(None, 4)

Total params: 123,568 (482.69 KB)

Trainable params: 123,568 (482.69 KB)

Non-trainable params: 0 (0.00 B)

Read Image Data

In [20]: `path = "/content/drive/MyDrive/DM3/updated"`


```
In [22]: batch_size = 16
dataset = tf.keras.preprocessing.image_dataset_from_directory(
    directory=path,
    labels='inferred',
    label_mode='categorical',
    batch_size=batch_size,
    validation_split=0.2,
    subset='training',
    seed=100
)


validation_dataset = tf.keras.preprocessing.image_dataset_from_directory(
    directory=path,
    labels='inferred',
    label_mode='categorical',
    batch_size=batch_size,
    validation_split=0.2,
    subset='validation',
    seed=100
)
```


```
Found 659 files belonging to 4 classes.
Using 528 files for training.
Found 659 files belonging to 4 classes.
Using 131 files for validation.
```


Train Model


```
In [23]: history = model.fit(dataset, validation_data=validation_dataset, epochs=20)
```


Epoch 1/20
33/33  92s 3s/step - accuracy: 0.2415 - loss: 56.8
145 - val_accuracy: 0.2366 - val_loss: 1.5235


Epoch 2/20
33/33  2s 57ms/step - accuracy: 0.3690 - loss: 1.3
270 - val_accuracy: 0.2290 - val_loss: 1.4999


Epoch 3/20
33/33  3s 70ms/step - accuracy: 0.4464 - loss: 1.2
150 - val_accuracy: 0.2901 - val_loss: 1.5305


Epoch 4/20
33/33  3s 93ms/step - accuracy: 0.4810 - loss: 1.1
148 - val_accuracy: 0.2901 - val_loss: 1.5962


Epoch 5/20
33/33  2s 58ms/step - accuracy: 0.5970 - loss: 0.9
940 - val_accuracy: 0.3053 - val_loss: 1.6106


Epoch 6/20
33/33  2s 66ms/step - accuracy: 0.6070 - loss: 0.9
415 - val_accuracy: 0.3282 - val_loss: 1.7379


Epoch 7/20
33/33  3s 67ms/step - accuracy: 0.6767 - loss: 0.8
280 - val_accuracy: 0.3511 - val_loss: 1.9491


Epoch 8/20
33/33  2s 58ms/step - accuracy: 0.6946 - loss: 0.7
848 - val_accuracy: 0.3359 - val_loss: 1.8639


Epoch 9/20
33/33  6s 154ms/step - accuracy: 0.7068 - loss: 0.
7332 - val_accuracy: 0.3053 - val_loss: 1.7994


Epoch 10/20
33/33  2s 61ms/step - accuracy: 0.7524 - loss: 0.7
232 - val_accuracy: 0.3740 - val_loss: 1.8349

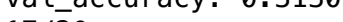
Epoch 11/20
33/33  2s 59ms/step - accuracy: 0.7411 - loss: 0.6
659 - val_accuracy: 0.3359 - val_loss: 1.8338

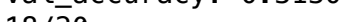
Epoch 12/20
33/33  2s 58ms/step - accuracy: 0.7812 - loss: 0.6
360 - val_accuracy: 0.3511 - val_loss: 1.8437


Epoch 13/20
33/33  2s 67ms/step - accuracy: 0.7704 - loss: 0.6
476 - val_accuracy: 0.3130 - val_loss: 1.8163


Epoch 14/20
33/33  2s 64ms/step - accuracy: 0.7901 - loss: 0.6
039 - val_accuracy: 0.3359 - val_loss: 1.8785

Epoch 15/20
33/33  3s 90ms/step - accuracy: 0.7931 - loss: 0.6
532 - val_accuracy: 0.3588 - val_loss: 2.1014

Epoch 16/20
33/33  4s 66ms/step - accuracy: 0.7755 - loss: 0.6
637 - val_accuracy: 0.3130 - val_loss: 1.8944

Epoch 17/20
33/33  2s 59ms/step - accuracy: 0.8109 - loss: 0.5
134 - val_accuracy: 0.3130 - val_loss: 1.9139

Epoch 18/20
33/33  2s 58ms/step - accuracy: 0.8044 - loss: 0.4
857 - val_accuracy: 0.3206 - val_loss: 2.1014

Epoch 19/20
33/33  3s 73ms/step - accuracy: 0.8146 - loss: 0.4
528 - val_accuracy: 0.3130 - val_loss: 2.4931

Epoch 20/20


33/33 ————— **3s 96ms/step** – accuracy: 0.8378 – loss: 0.4707 – val_accuracy: 0.3130 – val_loss: 2.3460

2 (a)

```
In [24]: training_accuracy = history.history['accuracy']
validation_accuracy = history.history['val_accuracy']

plt.plot(training_accuracy, label='Training Accuracy')
plt.plot(validation_accuracy, label='Validation Accuracy')
plt.title('Accuracy')
plt.xlabel('Epoch')
plt.ylabel('Accuracy')
plt.legend()

plt.tight_layout()
plt.show()
```



2 (b). Experiment with nodes changes to 4 and 16

```
In [25]: new_model1 = Sequential([
    Conv2D(8, (3, 3), activation='relu', input_shape=(256, 256, 3)),
    MaxPooling2D((2, 2)),
    Conv2D(8, (3, 3), activation='relu'),
    MaxPooling2D((2, 2)),
    Flatten(),
    Dense(8, activation='relu'),
    Dense(4, activation='softmax')
])

new_model2 = Sequential([
    Conv2D(8, (3, 3), activation='relu', input_shape=(256, 256, 3)),
    MaxPooling2D((2, 2)),
    Conv2D(16, (3, 3), activation='relu'),
    MaxPooling2D((2, 2)),
    Flatten(),
    Dense(8, activation='relu'),
    Dense(4, activation='softmax')
])
```

Compile new models

```
In [26]: new_model1.compile(optimizer='adam',
    loss='categorical_crossentropy',
    metrics=[tf.keras.metrics.CategoricalAccuracy(name='accuracy')])
new_model2.compile(optimizer='adam',
    loss='categorical_crossentropy',
    metrics=[tf.keras.metrics.CategoricalAccuracy(name='accuracy')])
```

```
In [27]: new_model1.summary()
```

Model: "sequential_1"

Layer (type)	Output Shape
conv2d_2 (Conv2D)	(None, 254, 254, 8)
max_pooling2d_2 (MaxPooling2D)	(None, 127, 127, 8)
conv2d_3 (Conv2D)	(None, 125, 125, 8)
max_pooling2d_3 (MaxPooling2D)	(None, 62, 62, 8)
flatten_1 (Flatten)	(None, 30752)
dense_2 (Dense)	(None, 8)
dense_3 (Dense)	(None, 4)

Total params: 246,868 (964.33 KB)

Trainable params: 246,868 (964.33 KB)

Non-trainable params: 0 (0.00 B)

```
In [28]: new_model2.summary()
```

Model: "sequential_2"

Layer (type)	Output Shape
conv2d_4 (Conv2D)	(None, 254, 254, 8)
max_pooling2d_4 (MaxPooling2D)	(None, 127, 127, 8)
conv2d_5 (Conv2D)	(None, 125, 125, 16)
max_pooling2d_5 (MaxPooling2D)	(None, 62, 62, 16)
flatten_2 (Flatten)	(None, 61504)
dense_4 (Dense)	(None, 8)
dense_5 (Dense)	(None, 4)


Total params: 493,468 (1.88 MB)


Trainable params: 493,468 (1.88 MB)


Non-trainable params: 0 (0.00 B)


Train new model 1


```
In [29]: new_history1 = new_model1.fit(dataset,validation_data=validation_data,epochs=20)
```


Epoch 1/20
33/33  7s 119ms/step - accuracy: 0.2305 - loss: 2
9.1985 - val_accuracy: 0.1985 - val_loss: 1.3867


Epoch 2/20
33/33  2s 58ms/step - accuracy: 0.2787 - loss: 1.3
851 - val_accuracy: 0.1985 - val_loss: 1.3872


Epoch 3/20
33/33  2s 59ms/step - accuracy: 0.3072 - loss: 1.3
845 - val_accuracy: 0.1985 - val_loss: 1.3876


Epoch 4/20
33/33  2s 58ms/step - accuracy: 0.2866 - loss: 1.3
839 - val_accuracy: 0.1985 - val_loss: 1.3882


Epoch 5/20
33/33  3s 67ms/step - accuracy: 0.2969 - loss: 1.3
833 - val_accuracy: 0.1985 - val_loss: 1.3889


Epoch 6/20
33/33  3s 90ms/step - accuracy: 0.2868 - loss: 1.3
833 - val_accuracy: 0.1985 - val_loss: 1.3891


Epoch 7/20
33/33  2s 67ms/step - accuracy: 0.2728 - loss: 1.3
846 - val_accuracy: 0.1985 - val_loss: 1.3896


Epoch 8/20
33/33  2s 58ms/step - accuracy: 0.2804 - loss: 1.3
829 - val_accuracy: 0.1985 - val_loss: 1.3902


Epoch 9/20
33/33  3s 58ms/step - accuracy: 0.2996 - loss: 1.3
803 - val_accuracy: 0.1985 - val_loss: 1.3908


Epoch 10/20
33/33  3s 58ms/step - accuracy: 0.2972 - loss: 1.3
809 - val_accuracy: 0.1985 - val_loss: 1.3912

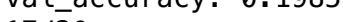
Epoch 11/20
33/33  3s 85ms/step - accuracy: 0.2816 - loss: 1.3
814 - val_accuracy: 0.1985 - val_loss: 1.3916

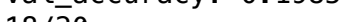
Epoch 12/20
33/33  3s 81ms/step - accuracy: 0.2927 - loss: 1.3
788 - val_accuracy: 0.1985 - val_loss: 1.3920


Epoch 13/20
33/33  4s 61ms/step - accuracy: 0.2894 - loss: 1.3
816 - val_accuracy: 0.1985 - val_loss: 1.3925


Epoch 14/20
33/33  2s 63ms/step - accuracy: 0.2929 - loss: 1.3
810 - val_accuracy: 0.1985 - val_loss: 1.3928

Epoch 15/20
33/33  3s 70ms/step - accuracy: 0.2928 - loss: 1.3
802 - val_accuracy: 0.1985 - val_loss: 1.3932


Epoch 16/20
33/33  3s 89ms/step - accuracy: 0.2744 - loss: 1.3
831 - val_accuracy: 0.1985 - val_loss: 1.3935

Epoch 17/20
33/33  3s 81ms/step - accuracy: 0.2871 - loss: 1.3
821 - val_accuracy: 0.1985 - val_loss: 1.3937

Epoch 18/20
33/33  2s 68ms/step - accuracy: 0.2869 - loss: 1.3
838 - val_accuracy: 0.1985 - val_loss: 1.3943


Epoch 19/20
33/33  2s 65ms/step - accuracy: 0.2756 - loss: 1.3
829 - val_accuracy: 0.1985 - val_loss: 1.3944


Epoch 20/20


33/33  **2s** 61ms/step – accuracy: 0.2776 – loss: 1.3842 – val_accuracy: 0.1985 – val_loss: 1.3947


Train New Model 2


```
In [30]: new_history2 = new_model2.fit(dataset,validation_data=validation_data,epochs=20)
```


Epoch 1/20
33/33  6s 98ms/step - accuracy: 0.3036 - loss: 20.8716 - val_accuracy: 0.2977 - val_loss: 1.3860


Epoch 2/20
33/33  4s 102ms/step - accuracy: 0.2478 - loss: 1.3862 - val_accuracy: 0.2977 - val_loss: 1.3861


Epoch 3/20
33/33  4s 110ms/step - accuracy: 0.2472 - loss: 1.3858 - val_accuracy: 0.2977 - val_loss: 1.3865


Epoch 4/20
33/33  2s 70ms/step - accuracy: 0.2636 - loss: 1.3853 - val_accuracy: 0.2977 - val_loss: 1.3867


Epoch 5/20
33/33  2s 62ms/step - accuracy: 0.2317 - loss: 1.3857 - val_accuracy: 0.1985 - val_loss: 1.3873


Epoch 6/20
33/33  2s 63ms/step - accuracy: 0.2860 - loss: 1.3840 - val_accuracy: 0.1985 - val_loss: 1.3877


Epoch 7/20
33/33  2s 69ms/step - accuracy: 0.2882 - loss: 1.3826 - val_accuracy: 0.1985 - val_loss: 1.3882


Epoch 8/20
33/33  2s 75ms/step - accuracy: 0.2933 - loss: 1.3835 - val_accuracy: 0.1985 - val_loss: 1.3888


Epoch 9/20
33/33  3s 90ms/step - accuracy: 0.2663 - loss: 1.3858 - val_accuracy: 0.1985 - val_loss: 1.3891


Epoch 10/20
33/33  4s 62ms/step - accuracy: 0.2762 - loss: 1.3831 - val_accuracy: 0.1985 - val_loss: 1.3897

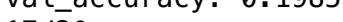
Epoch 11/20
33/33  2s 61ms/step - accuracy: 0.2883 - loss: 1.3812 - val_accuracy: 0.1985 - val_loss: 1.3902

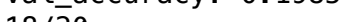
Epoch 12/20
33/33  2s 61ms/step - accuracy: 0.2666 - loss: 1.3837 - val_accuracy: 0.1985 - val_loss: 1.3907


Epoch 13/20
33/33  3s 78ms/step - accuracy: 0.2605 - loss: 1.3847 - val_accuracy: 0.1985 - val_loss: 1.3908


Epoch 14/20
33/33  3s 88ms/step - accuracy: 0.2706 - loss: 1.3821 - val_accuracy: 0.1985 - val_loss: 1.3915

Epoch 15/20
33/33  2s 60ms/step - accuracy: 0.2976 - loss: 1.3786 - val_accuracy: 0.1985 - val_loss: 1.3918

Epoch 16/20
33/33  2s 69ms/step - accuracy: 0.2799 - loss: 1.3826 - val_accuracy: 0.1985 - val_loss: 1.3923

Epoch 17/20
33/33  2s 67ms/step - accuracy: 0.2816 - loss: 1.3836 - val_accuracy: 0.1985 - val_loss: 1.3927

Epoch 18/20
33/33  2s 69ms/step - accuracy: 0.2697 - loss: 1.3846 - val_accuracy: 0.1985 - val_loss: 1.3929

Epoch 19/20
33/33  3s 82ms/step - accuracy: 0.2793 - loss: 1.3815 - val_accuracy: 0.1985 - val_loss: 1.3932

Epoch 20/20

33/33 ————— 3s 86ms/step – accuracy: 0.2881 – loss: 1.3811 – val_accuracy: 0.1985 – val_loss: 1.3935

2 (c)


```
In [31]: training_accuracy1 = new_history1.history['accuracy']
validation_accuracy1 = new_history1.history['val_accuracy']

training_accuracy2 = new_history2.history['accuracy']
validation_accuracy2 = new_history2.history['val_accuracy']

fig, axes = plt.subplots(1, 2, figsize=(15, 6))
axes[0].plot(training_accuracy1, label='Training Accuracy')
axes[0].plot(validation_accuracy1, label='Validation Accuracy')
axes[0].set_title('Accuracy')
axes[0].set_xlabel('Epoch')
axes[0].set_ylabel('Accuracy')

axes[1].plot(training_accuracy2, label='Training Accuracy')
axes[1].plot(validation_accuracy2, label='Validation Accuracy')
axes[1].set_title('Accuracy')
axes[1].set_xlabel('Epoch')
axes[1].set_ylabel('Accuracy')
plt.legend()

plt.show()
```



2 (d)

1. The initial model appears to suffer from overfitting. While the training accuracy increased to 84%, the validation accuracy remained below 25%.
2. In the experiment model 1 and model 2 suffers from underfitting, as persistently low accuracies across both training and validation sets.

3

```
In [33]: def load_data(file_path):  
         with open(file_path, 'r') as f:  
             data = [json.loads(line) for line in f]  
         return data
```

```
In [34]: train_data = load_data("/content/drive/MyDrive/DM3/student_23/train.js  
on")  
val_data = load_data("/content/drive/MyDrive/DM3/student_23/validatio  
n.json")  
test_data = load_data("/content/drive/MyDrive/DM3/student_23/test.js  
on")
```

```
In [36]: dataset = DatasetDict({  
         "train": Dataset.from_list(train_data),  
         "validation": Dataset.from_list(val_data),  
         "test": Dataset.from_list(test_data)  
})  
  
dataset.save_to_disk("sem_eval_dataset")
```

```
In [37]: labels = [label for label in dataset['train'].features.keys() if label  
not in ['ID', 'Tweet']]  
id2label = {idx:label for idx, label in enumerate(labels)}  
label2id = {label:idx for idx, label in enumerate(labels)}  
labels
```

```
Out[37]: ['anger',  
          'anticipation',  
          'disgust',  
          'fear',  
          'joy',  
          'love',  
          'optimism',  
          'pessimism',  
          'sadness',  
          'surprise',  
          'trust']
```

```
In [38]: from transformers import AutoTokenizer
import numpy as np

tokenizer = AutoTokenizer.from_pretrained("bert-base-uncased")

def preprocess_data(examples):
    text = examples["Tweet"]
    encoding = tokenizer(text, padding="max_length", truncation=True, max_length=128)
    labels_batch = {k: examples[k] for k in examples.keys() if k in labels}
    labels_matrix = np.zeros((len(text), len(labels)))
    for idx, label in enumerate(labels):
        labels_matrix[:, idx] = labels_batch[label]

    encoding["labels"] = labels_matrix.tolist()

    return encoding
```

```
In [39]: encoded_dataset = dataset.map(preprocess_data, batched=True, remove_columns=dataset['train'].column_names)
```

```
In [40]: encoded_dataset.set_format("torch")
```

```
In [41]: from transformers import AutoModelForSequenceClassification

model = AutoModelForSequenceClassification.from_pretrained("bert-base-uncased",
                                                         problem_type=
e="multi_label_classification",
                                                         num_labels=
len(labels),
                                                         id2label=id
2label,
                                                         label2id=label2id)
```

Some weights of BertForSequenceClassification were not initialized from the model checkpoint at bert-base-uncased and are newly initialized: ['classifier.bias', 'classifier.weight']
You should probably TRAIN this model on a down-stream task to be able to use it for predictions and inference.

```
In [42]: batch_size = 8
metric_name = "f1"
```

In [43]: `from transformers import TrainingArguments, Trainer`

```
args = TrainingArguments(
    f"bert-finetuned-sem_eval-english",
    evaluation_strategy = "epoch",
    save_strategy = "epoch",
    learning_rate=2e-5,
    per_device_train_batch_size=batch_size,
    per_device_eval_batch_size=batch_size,
    num_train_epochs=5,
    weight_decay=0.01,
    load_best_model_at_end=True,
    metric_for_best_model=metric_name,
)
```

In [44]: `from sklearn.metrics import f1_score, roc_auc_score, accuracy_score`
`from transformers import EvalPrediction`
`import torch`

```
def multi_label_metrics(predictions, labels, threshold=0.5):
    sigmoid = torch.nn.Sigmoid()
    probs = sigmoid(torch.Tensor(predictions))
    y_pred = np.zeros(probs.shape)
    y_pred[np.where(probs >= threshold)] = 1
    y_true = labels
    f1_micro_average = f1_score(y_true=y_true, y_pred=y_pred, average
                                = 'micro')
    roc_auc = roc_auc_score(y_true, y_pred, average = 'micro')
    accuracy = accuracy_score(y_true, y_pred)
    metrics = {'f1': f1_micro_average,
               'roc_auc': roc_auc,
               'accuracy': accuracy}
    return metrics

def compute_metrics(p: EvalPrediction):
    preds = p.predictions[0] if isinstance(p.predictions,
                                           tuple) else p.predictions
    result = multi_label_metrics(
        predictions=preds,
        labels=p.label_ids)
    return result
```

In [45]: `outputs = model(input_ids=encoded_dataset['train']['input_ids'][0].unsqueeze(0), labels=encoded_dataset['train'][0]['labels'].unsqueeze(0))`
`outputs`

We strongly recommend passing in an `attention_mask` since your input_ids may be padded. See <https://huggingface.co/docs/transformers/troubleshooting#incorrect-output-when-padding-tokens-arent-masked>.

Out [45]: `SequenceClassifierOutput(loss=tensor(0.6401, grad_fn=<BinaryCrossEntropyWithLogitsBackward0>), logits=tensor([[0.7788, -0.0009, 0.4671, 0.0596, -0.2869, -0.1744, -0.3475, 0.1844, 0.2262, -0.1921, 0.0752]], grad_fn=<AddmmBackward0>), hidden_states=None, attentions=None)`

```
In [46]: trainer = Trainer(  
    model,  
    args,  
    train_dataset=encoded_dataset["train"],  
    eval_dataset=encoded_dataset["validation"],  
    tokenizer=tokenizer,  
    compute_metrics=compute_metrics  
)
```

In [47]: `trainer.train()`

wandb: **WARNING** The `run_name` is currently set to the same value as `TrainingArguments.output_dir`. If this was not intended, please specify a different run name by setting the `TrainingArguments.run_name` parameter.

wandb: Using wandb-core as the SDK backend. Please refer to <https://wandb.me/wandb-core> for more information.

wandb: Logging into wandb.ai. (Learn how to deploy a W&B server locally: <https://wandb.me/wandb-server>)

wandb: You can find your API key in your browser here: <https://wandb.ai/authorize>

wandb: Paste an API key from your profile and hit enter, or press ctrl+c to quit:

.....

wandb: **ERROR** API key must be 40 characters long, yours was 25

wandb: Appending key for api.wandb.ai to your netrc file: /root/.netrc

Tracking run with wandb version 0.18.7

Run data is saved locally in /content/wandb/run-20241129_220424-u04sy97i

Syncing run **bert-finetuned-sem eval-english** (<https://wandb.ai/patels78-rowan-university/huggingface/runs/u04sy97i>) to **Weights & Biases** (<https://wandb.ai/patels78-rowan-university/huggingface>) ([docs \(https://wandb.me/developer-guide\)](https://wandb.me/developer-guide))

View project at <https://wandb.ai/patels78-rowan-university/huggingface> (<https://wandb.ai/patels78-rowan-university/huggingface>)

View run at <https://wandb.ai/patels78-rowan-university/huggingface/runs/u04sy97i> (<https://wandb.ai/patels78-rowan-university/huggingface/runs/u04sy97i>)

[1875/1875 07:33, Epoch 5/5]

Epoch	Training Loss	Validation Loss	F1	Roc Auc	Accuracy
1	No log	0.354214	0.617700	0.736169	0.192500
2	0.407500	0.327928	0.643457	0.750446	0.220000
3	0.296000	0.322414	0.681406	0.778505	0.247500
4	0.246800	0.320954	0.663578	0.765654	0.220000
5	0.246800	0.320341	0.692521	0.788160	0.250000

Out [47]: TrainOutput(global_step=1875, training_loss=0.29769027099609374, metrics={'train_runtime': 519.0228, 'train_samples_per_second': 28.9, 'train_steps_per_second': 3.613, 'total_flos': 986746187520000.0, 'train_loss': 0.29769027099609374, 'epoch': 5.0})


```
In [48]: test_results = trainer.evaluate(encoded_dataset["test"])
print("Test Results:", test_results)
```

[188/188 00:10]

```
Test Results: {'eval_loss': 0.31512686610221863, 'eval_f1': 0.68299319
72789116, 'eval_roc_auc': 0.7830501218682399, 'eval_accuracy': 0.27133
3333333333, 'eval_runtime': 10.7438, 'eval_samples_per_second': 139.6
16, 'eval_steps_per_second': 17.499, 'epoch': 5.0}
```

```
In [52]: train_losses = [log["loss"] for log in trainer.state.log_history if "l
oss" in log]
train_losses += [train_losses[-1]]*3
val_losses = [log["eval_loss"] for log in trainer.state.log_history if
"eval_loss" in log]

epochs = range(1, len(train_losses) + 1)
plt.figure(figsize=(10, 6))
plt.plot(epochs, train_losses, label="Training Loss", marker="o")
plt.plot(epochs, val_losses, label="Validation Loss", marker="o")
plt.xlabel("Epochs")
plt.ylabel("Loss")
plt.title("Learning Curves")
plt.legend()
plt.grid()
plt.show()
```



```
In [50]: from sklearn.metrics import accuracy_score
import numpy as np

predictions = trainer.predict(encoded_dataset["test"])
preds_logits = predictions.predictions
true_labels = predictions.label_ids

sigmoid = torch.nn.Sigmoid()
probs = sigmoid(torch.Tensor(preds_logits)).numpy()

threshold = 0.5
binary_preds = (probs >= threshold).astype(int)

def exact_match_accuracy(predictions, labels):
    return np.mean(np.all(predictions == labels, axis=1))

def at_least_one_match_accuracy(predictions, labels):
    acc=0
    for i,j in zip(binary_preds,true_labels):
        for(a,b) in zip(i,j):
            if((a==b) and a == 1):
                acc +=1
                break
    return acc/len(predictions)

exact_match_acc = exact_match_accuracy(binary_preds, true_labels)
at_least_one_match_acc = at_least_one_match_accuracy(binary_preds, true_labels)

print(f"Exact Match Accuracy: {exact_match_acc * 100:.2f}%")
print(f"At Least One Match Accuracy: {at_least_one_match_acc * 100:.2f}%")
```

Exact Match Accuracy: 27.13%
At Least One Match Accuracy: 85.87%